



## Department of Electrical Engineering

**Faculty Member:** LE Munadi Sial  
**Semester:**

**Date:**  
**Group:**

### CS471 Machine Learning

#### **Lab 7: Linear Regression II – Train-Validation Split and Regularization**

Name	Reg. No	PLO4 - CLO4	PLO4 - CLO4	PLO5 - CLO5	PLO8 - CLO6	PLO9 - CLO7
		Viva /Quiz / Lab Performance	Analysis of data in Lab Report	Modern Tool Usage	Ethics	Individual and Team Work
		5 Marks	5 Marks	5 Marks	5 Marks	5 Marks
Muhammad Anser Sohaib	367628					
Muhammad Taqi	372071					
Muhammad Talha	370860					
Zuha Fatima	385647					
Hashaam Zafar	395508					



## **Introduction**

This laboratory exercise will extend the python implementation of linear regression performed in the previous lab. Linear regression is a basic supervised learning technique in which parameters are trained on a dataset to fit a model that best approximates that dataset. The problem with using simple linear regression is that the trained models can overfit the dataset at which point regularization must be used to prevent overfitting. This lab will focus on integrating the regularization concept into the gradient descent algorithm.

## **Objectives**

The following are the main objectives of this lab:

- Extract and prepare the training and cross-validation datasets
- Use feature scaling to ensure uniformity among the feature columns
- Implement cost function on both training and cross-validation datasets
- Implement gradient descent algorithm
- Plot the training and cross-validation losses
- Use L2 regularization to counter overfitting

## **Lab Conduct**

- Respect faculty and peers through speech and actions
- The lab faculty will be available to assist the students. In case some aspect of the lab experiment is not understood, the students are advised to seek help from the faculty.
- In the tasks, there are commented lines such as `#YOUR CODE STARTS HERE#` where you have to provide the code. You must put the code/screenshot/plot between the `#START` and `#END` parts of these commented lines. Do NOT remove the commented lines.



- Use the tab key to provide the indentation in python.
- When you provide the code in the report, keep the font size at 12

## Theory

Linear Regression is a very basic supervised learning technique. To calculate the loss in each training example, the difference between a hypothesis and the label ( $y$ ) is calculated. The hypothesis is a linear equation of the features ( $x$ ) in the dataset with the coefficients acting as the weight parameters. These weight parameters are initialized to random values at the start but are then trained over time to learn the model. The cost function is used to calculate the error between the predicted  $\hat{y}$  and the actual  $y$ .

A major problem in the training is that the weights that are trained may fit the model for only the data it is given. This means that the model will not generalize to examples outside the dataset and is referred to as “overfitting”. Such overfitting makes the machine learning implementation very impractical for real-life applications where data has high variation. To prevent overfitting of the model, a modification in the cost function and gradient descent is implemented. This modification is called regularization and is itself controlled by a hyperparameter ( $\lambda$ ).

A brief summary of the relevant keywords and functions in python is provided below:

<b>print()</b>	output text on console
<b>input()</b>	get input from user on console
<b>range()</b>	create a sequence of numbers
<b>len()</b>	gives the number of characters in a string
<b>if</b>	contains code that executes depending on a logical condition
<b>else</b>	connects with <b>if</b> and <b>elif</b> , executes when conditions are not met
<b>elif</b>	equivalent to <b>else if</b>
<b>while</b>	loops code as long as a condition is true
<b>for</b>	loops code through a sequence of items in an iterable object
<b>break</b>	exit loop immediately
<b>continue</b>	jump to the next iteration of the loop



**def**                      used to define a function

## Lab Task 1 - Dataset Preparation, Feature Scaling \_\_\_\_\_

You have been provided with a dataset containing several feature columns. You will need to select any 3 of the feature columns to make your own dataset. The “Sale Price” is the label column that your model will predict. The dataset examples are to be divided into 2 separate portions: training and cross-validation datasets (choose from 80-20 to 70-30 ratios). Save the prepared datasets as CSV files. Next, load the datasets into your python program and store them as NumPy arrays ( $X_{\text{train}}$ ,  $y_{\text{train}}$ ,  $X_{\text{val}}$ ,  $y_{\text{val}}$ ). Next, use feature scaling to rescale the feature columns of both datasets so that their values range from 0 to 1. Finally, print both of the datasets (you need to show any 5 rows of the datasets).

### ### TASK 1 CODE STARTS HERE ###

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

dataset = pd.read_csv('ML_Lab7_Dataset.csv')
refined_dataset = dataset[["LotArea", "OverallQual", "YrSold", "SalePrice"]]

def normalize(df):
    scaler = StandardScaler()
    df = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)
    return df

refined_dataset = normalize(refined_dataset)

def split_dataset(dataset, test_ratio):
    test_size = len(dataset) * test_ratio
    X_train = dataset.take(range(0, int(len(dataset) - test_size)))
    X_test = dataset.take(range(int(len(dataset) - test_size), len(dataset)))
    y_train = X_train.pop("SalePrice")
```



```
y_test = X_test.pop("SalePrice")  
return X_train, X_test, y_train, y_test  
  
X_train, X_test, y_train, y_test = split_dataset(refined_dataset, 0.2)
```

### TASK 1 CODE ENDS HERE ###

### TASK 1 SCREENSHOT STARTS HERE ###

The screenshot shows two Jupyter Notebook cells. The first cell contains the code `X_train.head()` and displays a DataFrame with 5 rows and 4 columns: `LotArea`, `OverallQual`, and `YrSold`. The second cell contains the code `y_train.head()` and displays a Series with 5 rows and 1 column: `SalePrice`. The data type for the second series is `dtype: float64`.

	LotArea	OverallQual	YrSold
0	-0.207142	0.651479	0.138777
1	-0.091886	-0.071836	-0.614439
2	0.073480	0.651479	0.138777
3	-0.096897	0.651479	-1.367655
4	0.375148	1.374795	0.138777

	SalePrice
0	0.347273
1	0.007288
2	0.536154
3	-0.515281
4	0.869843

dtype: float64

### TASK 1 SCREENSHOT ENDS HERE ###



## Lab Task 2 - Cost Function with Regularization

---

For linear regression, you will implement the following hypothesis:

$$h(x) = w_0 + w_1x_1 + w_2x_2 + w_3x_3 + \dots$$

The  $w_j$  and  $b$  represent the weights while the  $x_j$  represents the  $j^{\text{th}}$  feature. The linear hypothesis  $h(x)$  is to be calculated for each training example and its difference with the label  $y$  of that training example will represent the loss. In this task, you will write a cost function that calculates the overall loss across a set of examples. This cost function will be useful to calculate the losses in both the training and cross-validation phases of the program.

`cost_function(X, y, lambd)`

The  $X$  and  $y$  are the features and labels of either the training or the cross-validation datasets. This is useful as it can be used for either the training examples or the cross-validation examples of the dataset. The *lambd* is the regularization parameter (Note that *lambda* is a keyword reserved in python). The function will calculate the losses to return the overall cost value. The cost function is given by:

$$J(w) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2 + \frac{1}{n} \lambda \sum_{j=1}^n (w_j)^2$$

The  $m$  is the number of the examples in the dataset and  $n$  is the total number of features (or non-bias weights) in the hypothesis. Write the code for the cost function and implement it for your training and cross-validation datasets to print out the cost. Provide the code and all relevant screenshots of the final output.

**### TASK 2 CODE STARTS HERE ###**



```
import numpy as np

def cost_with_regularization(X, y, lambd, w0, w1, w2, w3):

    m = len(X)
    predictions = w0 + X[:, 0] * w1 + X[:, 1] * w2 + X[:, 2] * w3

    # Compute the cost without regularization
    cost = (1 / (2 * m)) * np.sum((predictions - y) ** 2)

    # Add regularization term (note: w0 is excluded)
    reg_term = (lambd / (2 * m)) * (w1**2 + w2**2 + w3**2)
    total_cost = cost + reg_term

    return total_cost
print(cost_with_regularization(X_train, y_train, 0.01, 0.5, 0.01, 0.2, 0.001))
```

**### TASK 2 CODE ENDS HERE ###**

**### TASK 2 SCREENSHOT STARTS HERE ###**

Cost for random weights and regularization is: 0.479658616205666

**### TASK 2 SCREENSHOT ENDS HERE ###**

### Lab Task 3 –Gradient Descent with Regularization \_\_\_\_\_

In this task, you will write a function that uses gradient descent to update the weight parameters:

```
gradient_descent(X, y, alpha, lambd)
```





The *alpha* is the learning rate (hyperparameter 1) and *lambda* is the regularization parameter (hyperparameter 2). The gradient descent algorithm is given as follows:

$$dw_j = \frac{\partial J}{\partial w_j} = \frac{1}{m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})x_j^{(i)} + \frac{\lambda}{m} w_j$$

$$db = \frac{\partial J}{\partial b} = \frac{1}{m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})$$

$$w_j := w_j - \alpha \frac{\partial J}{\partial w_j}$$

$$b := b - \alpha \frac{\partial J}{\partial b}$$

For the submission, you will need to run the gradient descent algorithm once to update the weights. You will need to print the weights, training cost and validation cost both before and after the weight update. Provide the code and all relevant screenshots of the final output.

### ### TASK 3 CODE STARTS HERE ###

```
def gradient_descent_with_regularization(X, y, alpha, lambda, num_epochs):  
    w0 = np.random.randn() * 0.01  
    w1 = np.random.randn() * 0.01  
    w2 = np.random.randn() * 0.01  
    w3 = np.random.randn() * 0.01  
    for i in range(num_epochs):  
  
        predictions = (w0 + X[:, 0] * w1 + X[:, 1] * w2 + X[:, 2] * w3)
```





```
dw0 = 1 / len(X) * np.sum(predictions - y)
dw1 = (1 / len(X) * np.sum((predictions - y) * X[:, 0])) + (lambda /
len(X)) * w1
dw2 = (1 / len(X) * np.sum((predictions - y) * X[:, 1])) + (lambda /
len(X)) * w2
dw3 = (1 / len(X) * np.sum((predictions - y) * X[:, 2])) + (lambda /
len(X)) * w3

w0 -= alpha * dw0
w1 -= alpha * dw1
w2 -= alpha * dw2
w3 -= alpha * dw3

# Log the cost every 50 epochs for better tracking
if i % 2 == 0:
    cost = cost_with_regularization(X, y, lambda, w0, w1, w2, w3)
    print(f"Epoch {i}: Cost = {cost:.4f}") # Ensure cost is a scalar

return w0, w1, w2, w3 # Return the weights after training
```

**### TASK 3 CODE ENDS HERE ###**

**### TASK 3 SCREENSHOT STARTS HERE ###**

```
Epoch 0: Cost = 0.4654
Epoch 1: Cost = 0.4591
Final Training Cost: 0.4591
Final Testing Cost: 0.4992
```

**### TASK 3 SCREENSHOT ENDS HERE ###**

## Lab Task 4 – Training and Validation Program \_\_\_\_\_



In this task, you will use the functions from the previous two tasks to write a “main” function that performs the actual training and validation. Use the cost function and gradient descent function on the training examples to determine the training loss and update the weights respectively. Then, use the cost function on the cross-validation examples to determine the cross-validation loss. This single iteration over the entire dataset (both training and cross-validation) marks the completion of one epoch. You will need to perform the training and cross-validation over several epochs (the epoch number is another hyperparameter that must be chosen). Ensure that at the end of each epoch, the training and cross-validation losses are stored for plotting purposes. When the final epoch is performed, note down the trained parameters (weights and bias) and make plot of the training and cross-validation losses (y-axis) over the epochs (x-axis). Ensure that both of the losses appear on the same graph. You only need to show a single plot for this task. Provide the code (excluding function definitions) and all relevant screenshots of the final output.

### ### TASK 4 CODE STARTS HERE ###

```
import numpy as np

def cost_with_regularization(X, y, lambd, w0, w1, w2, w3):
    m = len(X)
    predictions = w0 + X[:, 0] * w1 + X[:, 1] * w2 + X[:, 2] * w3

    # Compute the cost without regularization
    cost = (1 / (2 * m)) * np.sum((predictions - y) ** 2)

    # Add regularization term (note: w0 is excluded)
    reg_term = (lambd / (2 * m)) * (w1**2 + w2**2 + w3**2)
    total_cost = cost + reg_term

    return total_cost

def gradient_descent_with_regularization(X, y, alpha, lambd, num_epochs):
    w0 = np.random.randn() * 0.01
    w1 = np.random.randn() * 0.01
    w2 = np.random.randn() * 0.01
```



```
w3 = np.random.randn() * 0.01
for i in range(num_epochs):

    predictions = (w0 + X[:, 0] * w1 + X[:, 1] * w2 + X[:, 2] * w3)

    dw0 = 1 / len(X) * np.sum(predictions - y)
    dw1 = (1 / len(X) * np.sum((predictions - y) * X[:, 0])) + (lambda /
len(X)) * w1
    dw2 = (1 / len(X) * np.sum((predictions - y) * X[:, 1])) + (lambda /
len(X)) * w2
    dw3 = (1 / len(X) * np.sum((predictions - y) * X[:, 2])) + (lambda /
len(X)) * w3

    w0 -= alpha * dw0
    w1 -= alpha * dw1
    w2 -= alpha * dw2
    w3 -= alpha * dw3

    # Log the cost every 50 epochs for better tracking
    if i % 30 == 0:
        cost = cost_with_regularization(X, y, lambda, w0, w1, w2, w3)
        print(f"Epoch {i}: Cost = {cost:.4f}") # Ensure cost is a scalar

    return w0, w1, w2, w3 # Return the weights after training

# Assuming X_train, y_train, X_test, y_test are defined before this
w0, w1, w2, w3 = gradient_descent_with_regularization(X_train, y_train, 0.01,
0.01, 1000)
train_cost = cost_with_regularization(X_train, y_train, 0.01, w0, w1, w2, w3)
test_cost = cost_with_regularization(X_test, y_test, 0.01, w0, w1, w2, w3)
print(f"Final Training Cost: {train_cost:.4f}")
print(f"Final Testing Cost: {test_cost:.4f}")
```

**### TASK 4 CODE ENDS HERE ###**

**### TASK 4 SCREENSHOT STARTS HERE ###**



```
⇒ Epoch 0: Cost = 0.4719  
Epoch 30: Cost = 0.3278  
Epoch 60: Cost = 0.2518  
Epoch 90: Cost = 0.2114  
Epoch 120: Cost = 0.1899  
Epoch 150: Cost = 0.1783  
Epoch 180: Cost = 0.1720  
Epoch 210: Cost = 0.1686  
Epoch 240: Cost = 0.1667  
Epoch 270: Cost = 0.1657  
Epoch 300: Cost = 0.1652  
Epoch 330: Cost = 0.1648  
Epoch 360: Cost = 0.1647  
Epoch 390: Cost = 0.1646  
Epoch 420: Cost = 0.1645  
Epoch 450: Cost = 0.1645  
Epoch 480: Cost = 0.1645  
Epoch 510: Cost = 0.1645  
Epoch 540: Cost = 0.1645  
Epoch 570: Cost = 0.1645  
Epoch 600: Cost = 0.1645  
Epoch 630: Cost = 0.1645  
Epoch 660: Cost = 0.1644  
Epoch 690: Cost = 0.1644  
Epoch 720: Cost = 0.1644  
Epoch 750: Cost = 0.1644  
Epoch 780: Cost = 0.1644  
Epoch 810: Cost = 0.1644  
Epoch 840: Cost = 0.1644  
Epoch 870: Cost = 0.1644  
Epoch 900: Cost = 0.1644  
Epoch 930: Cost = 0.1644  
Epoch 960: Cost = 0.1644  
Epoch 990: Cost = 0.1644  
Final Training Cost: 0.1644  
Final Testing Cost: 0.1962
```

### TASK 4 SCREENSHOT ENDS HERE ###

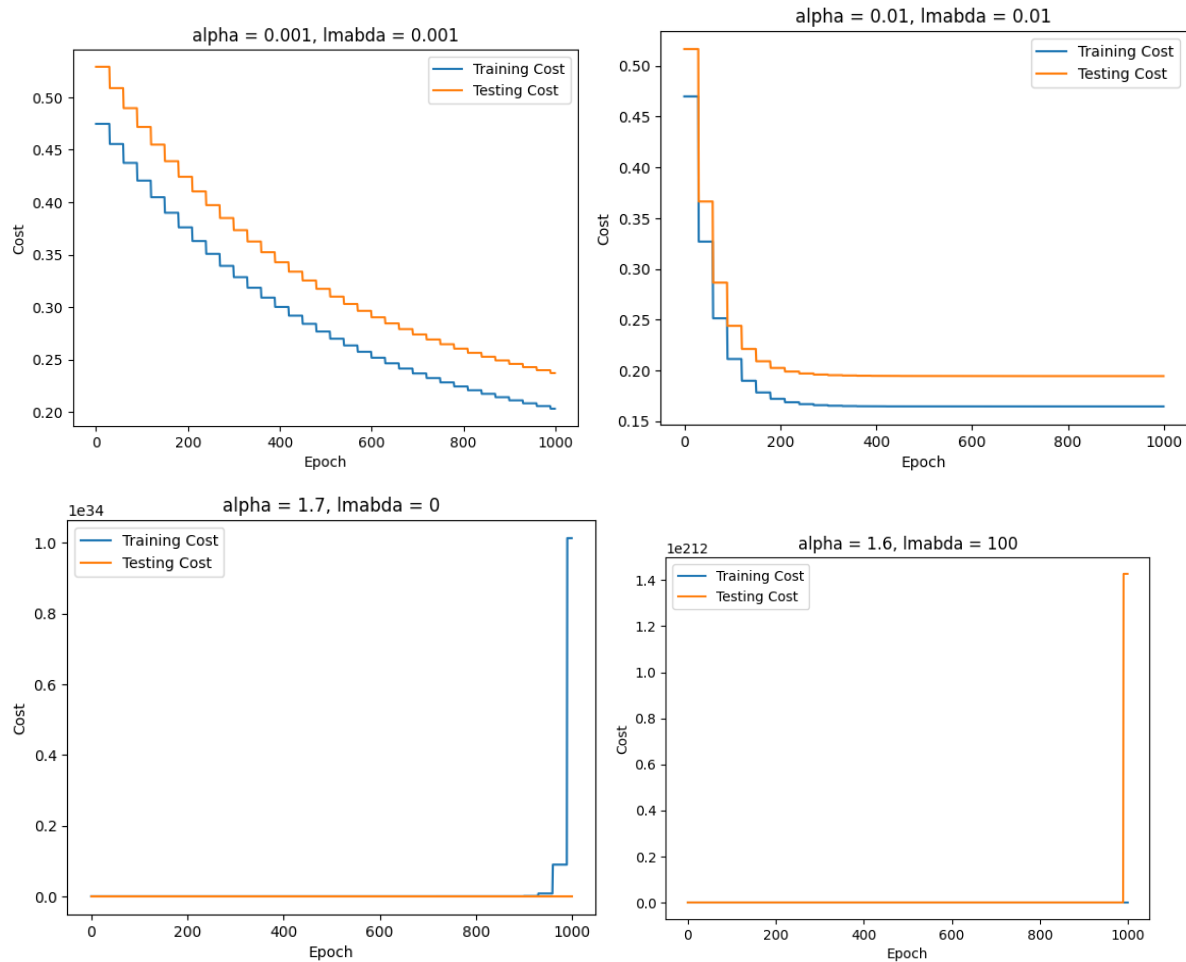
## Lab Task 5 – Tuning Alpha and Lambda

In this task, you will use your linear regression code from the previous task. Tune the alpha and lambda hyperparameters at different values to get several plots.



You need to get at least 6 plots. Mention the alpha and lambda values in the plot titles. Ensure all axes are labeled appropriately.

### ### TASK 5 PLOTS START HERE ###



### ### TASK 5 PLOTS END HERE ###